

# Implementing Graph Data Models in Neo4j

# Lesson Overview

Modules in this course:

1. Implementing Your First Model
2. Importing Data
3. Profiling Queries
4. Refactoring Graphs

# Implementing Graph Data Models in Neo4j



# In This Module You'll Learn ...

At the end of this module, you should be able to:

- Write Cypher code to implement a simple graph data model
- Confirm that the starter data is in the graph

# Model Domain

<https://www.bts.gov/browse-statistical-products-and-data>



The screenshot shows the homepage of the Bureau of Transportation Statistics (BTS) website. The header is dark blue with the United States Department of Transportation logo and name. The main navigation bar is also dark blue and contains the text "Bureau of Transportation Statistics" and a search bar labeled "Search BTS site". Below the navigation bar, there are five tabs: "Topics and Geography", "Statistical Products and Data", "National Transportation Library", "Newsroom", and "About BTS". The main content area is white and features the heading "Browse Statistical Products and Data". Below this heading, there is a paragraph: "Browse databases, surveys, reports and other statistical products created by BTS. Use BTS tools and apps to build your own statistical tables, graphics and maps." At the bottom of the page, there are three images: a laptop displaying a line graph, a smartphone displaying a map of the United States, and a printed report with a map and a table.

United States Department of Transportation

Ask-A-Librarian | A-Z Index

## Bureau of Transportation Statistics

Search BTS site

Topics and Geography | Statistical Products and Data | National Transportation Library | Newsroom | About BTS

Home

### Browse Statistical Products and Data

Browse databases, surveys, reports and other statistical products created by BTS. Use BTS tools and apps to build your own statistical tables, graphics and maps.

5



# Model Data

Origin		
<a href="#">OriginAirportID</a>	Origin Airport, Airport ID. An identification number assigned by US DOT to identify a unique airport. Use this field for airport analysis across a range of years because an airport can change its airport code and airport codes can be reused.	<a href="#">Analysis</a>
<a href="#">OriginAirportSeqID</a>	Origin Airport, Airport Sequence ID. An identification number assigned by US DOT to identify a unique airport at a given point of time. Airport attributes, such as airport name or coordinates, may change over time.	
<a href="#">OriginCityMarketID</a>	Origin Airport, City Market ID. City Market ID is an identification number assigned by US DOT to identify a city market. Use this field to consolidate airports serving the same city market.	<a href="#">Analysis</a>
<a href="#">Origin</a>	Origin Airport	<a href="#">Analysis</a>
<a href="#">OriginCityName</a>	Origin Airport, City Name	
<a href="#">OriginState</a>	Origin Airport, State Code	<a href="#">Analysis</a>
<a href="#">OriginStateFips</a>	Origin Airport, State Fips	<a href="#">Analysis</a>
<a href="#">OriginStateName</a>	Origin Airport, State Name	
<a href="#">OriginWac</a>	Origin Airport, World Area Code	<a href="#">Analysis</a>

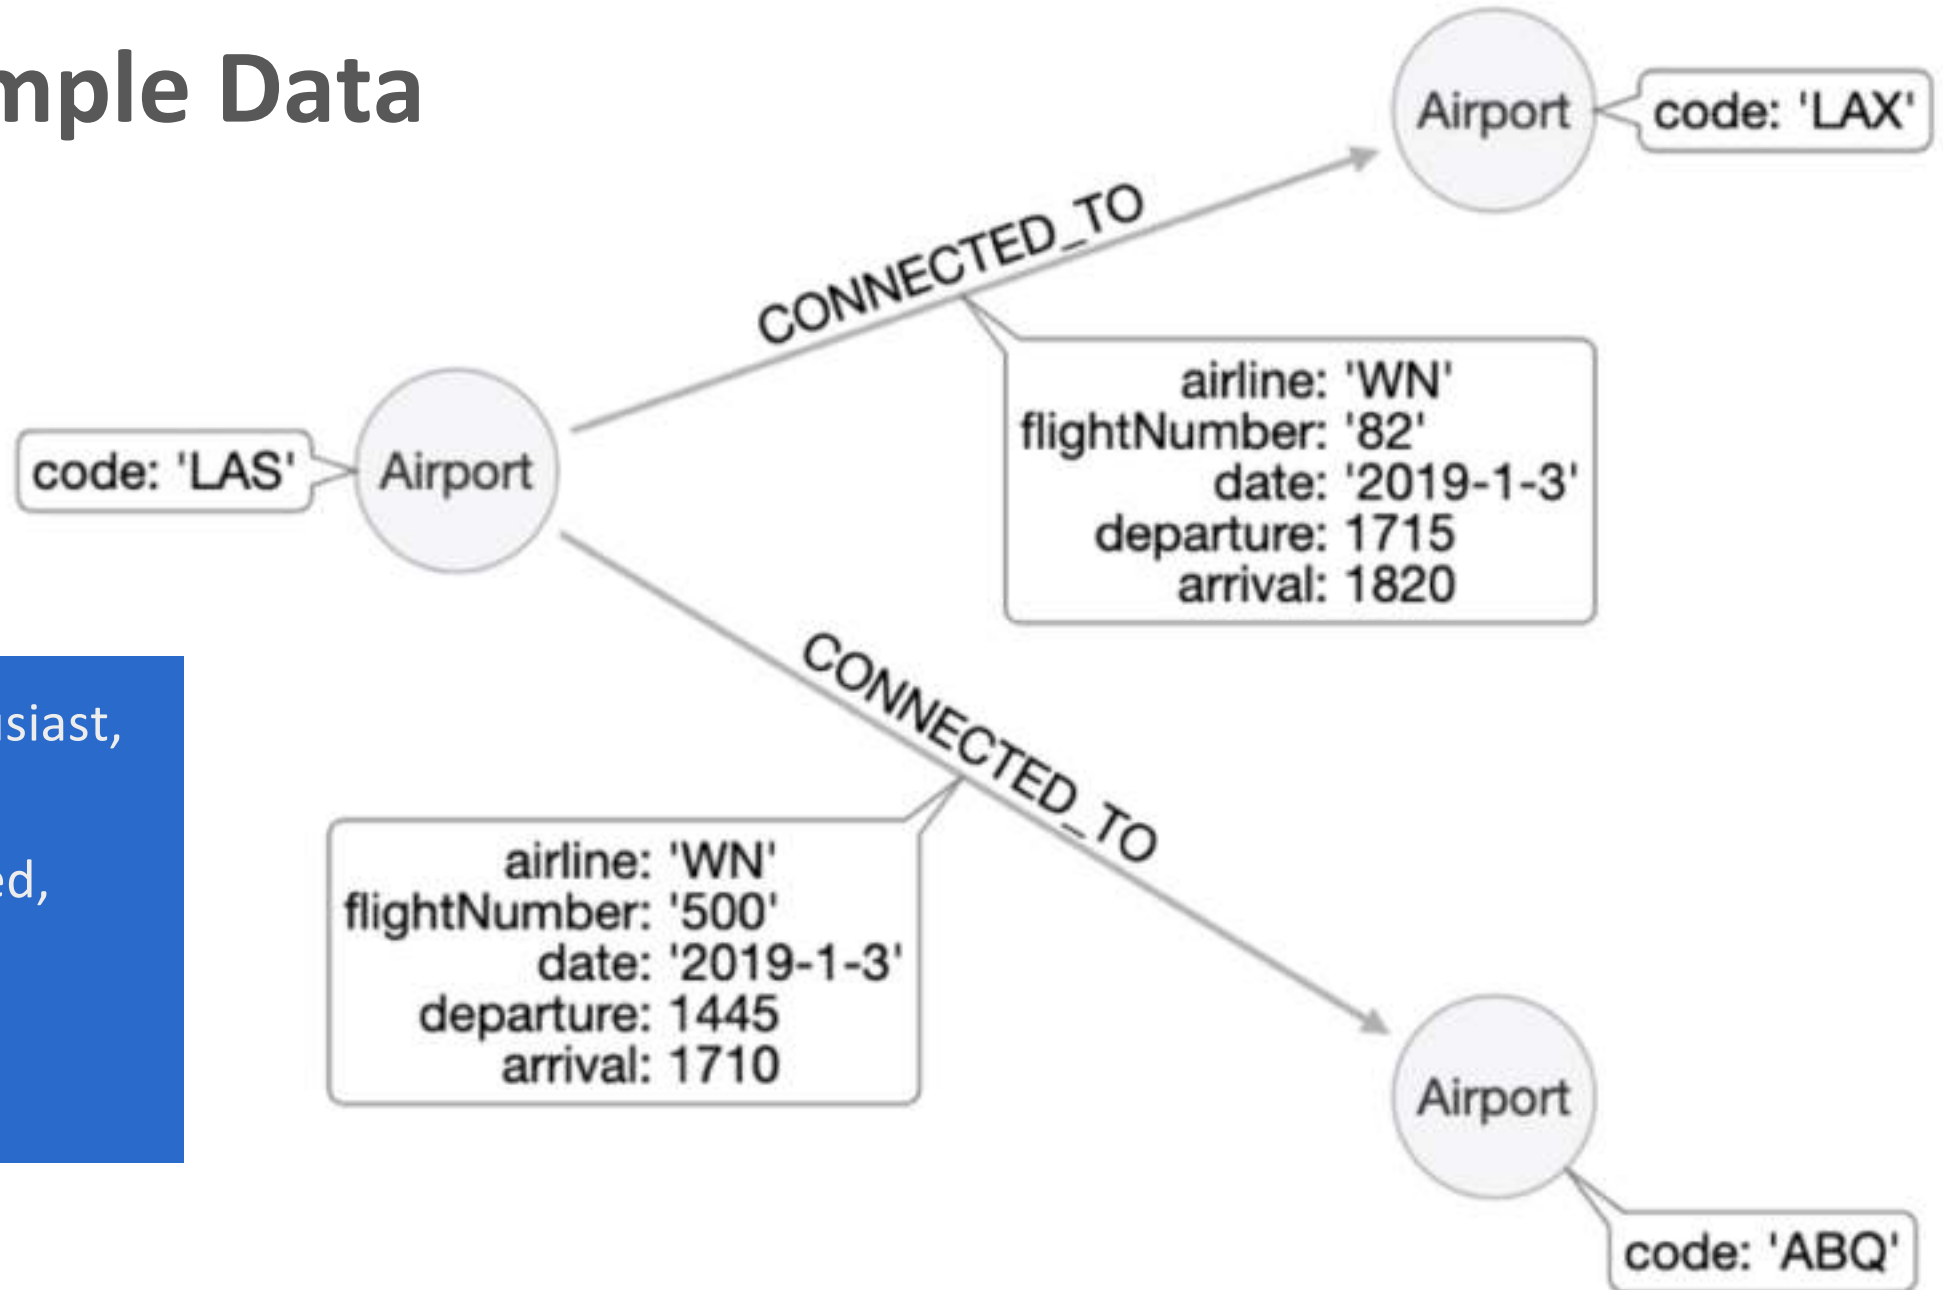
Destination		
<a href="#">DestAirportID</a>	Destination Airport, Airport ID. An identification number assigned by US DOT to identify a unique airport. Use this field for airport analysis across a range of years because an airport can change its airport code and airport codes can be reused.	<a href="#">Analysis</a>
<a href="#">DestAirportSeqID</a>	Destination Airport, Airport Sequence ID. An identification number assigned by US DOT to identify a unique airport at a given point of time. Airport attributes, such as airport name or coordinates, may change over time.	
<a href="#">DestCityMarketID</a>	Destination Airport, City Market ID. City Market ID is an identification number assigned by US DOT to identify a city market. Use this field to consolidate airports serving the same city market.	<a href="#">Analysis</a>
<a href="#">Dest</a>	Destination Airport	<a href="#">Analysis</a>
<a href="#">DestCityName</a>	Destination Airport, City Name	
<a href="#">DestState</a>	Destination Airport, State Code	<a href="#">Analysis</a>
<a href="#">DestStateFips</a>	Destination Airport, State Fips	<a href="#">Analysis</a>
<a href="#">DestStateName</a>	Destination Airport, State Name	
<a href="#">DestWac</a>	Destination Airport, World Area Code	<a href="#">Analysis</a>

Time Period		
<a href="#">Year</a>	Year	
<a href="#">Quarter</a>	Quarter (1-4)	<a href="#">Analysis</a>
<a href="#">Month</a>	Month	<a href="#">Analysis</a>
<a href="#">DayOfMonth</a>	Day of Month	
<a href="#">DayOfWeek</a>	Day of Week	<a href="#">Analysis</a>
<a href="#">FlightDate</a>	Flight Date (yyyymmdd)	

# Model Application Question

As an air travel enthusiast,  
I want to know how airports are connected,  
so I find the busiest airports.

# Model Sample Data



As an air travel enthusiast,  
I want to know how  
airports are connected,  
so I find the busiest  
airports.



# Exercise 1: Getting Started with the Airport Graph Data Model

Before you start this exercise you must:

1. Create a project in Neo4j Desktop, create a blank sandbox, or create a Neo4j Aura instance
2. If using Neo4j Desktop, create a local 4.x database in the project and start it
3. Open a Neo4j Browser window for the database

In the query edit pane of Neo4j Browser, execute the browser command:

```
:play 4.0-neo4j-modeling-exercises
```

and follow the instructions for Exercise 1.

Note: This exercise has 3 steps. Estimated time to complete: 15 minutes.



# Importing Data into the Graph

# In This Module You'll Learn ...

At the end of this module, you should be able to:

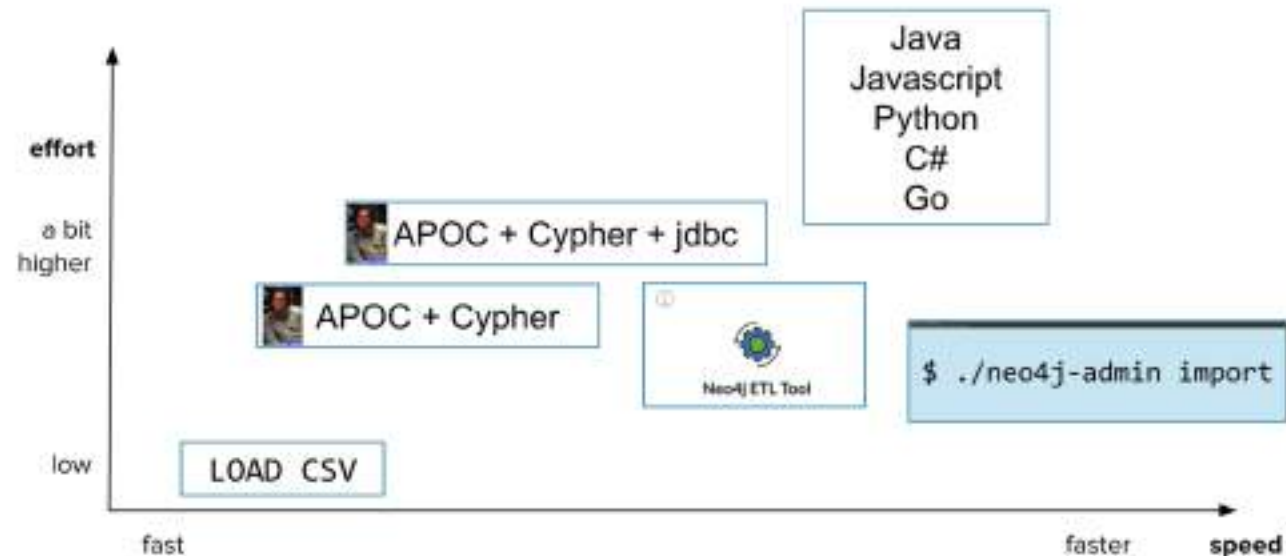
- Write Cypher code to import CSV data into a graph
- Confirm that the data has been loaded

# Options for Importing Data into a Graph

There are many options for importing data into Neo4j

The option selected for importing data depends on ...

- The **amount of data** to be imported
- The **available tools** and knowledge of those tools
- The **amount of time** allocated to perform the import



# Import Preparation

- Names of entities (node labels)
- Names of relationships
- Names of properties for nodes and relationships
- Constraints to be defined
- Indexes required
- The most important queries

# LOAD CSV Syntax

Simplified syntax for using LOAD CSV:

```
LOAD CSV      // load csv data
WITH HEADERS  // optionally use first header row as keys in "row" map
FROM "url"     // file:/// file relative to $NEO4J_HOME/import or http://
AS row        // return each row of the CSV as list of strings or map
// ... rest of the Cypher statement ...
```

Use **LOAD CSV** for CSV files that contain fewer than 100k lines



# Inspecting CSV File Data on a Network

```
1 LOAD CSV WITH HEADERS FROM 'https://r.neo4j.com/flights_2019_1k' AS row
2 RETURN row
3 LIMIT 5
```

row

Table

Text

Code

```
{
  "FlightNum": "335",
  "Origin": "IAD",
  "LateAircraftDelay": "NA",
  "NASDelay": "NA",
  "ArrTime": "2211",
  "AirTime": "116",
  "DepTime": "2003",
  "Month": "1",
  "CRSElapsedTime": "150",
  "DayofMonth": "3",
  "Distance": "810",
  "CRSDepTime": "1955",
  "SecurityDelay": "NA",
  "DayOfWeek": "4"
```

Started streaming 5 records after 1 ms and completed after 954 ms.

# Creating Nodes and Relationships

**LOAD CSV** command reads rows of data from a CSV file

- It then creates nodes and relationships in the graph

```
LOAD CSV WITH HEADERS FROM 'https://r.neo4j.com/flights_2019_1k' AS row
MERGE (origin:Airport {code: row.Origin})
MERGE (destination:Airport {code: row.Dest})
MERGE (origin)-[connection:CONNECTED_TO {
    airline: row.UniqueCarrier,
    flightNumber: row.FlightNum,
    date: toInteger(row.Year) + '-' + toInteger(row.Month) + '-' +
        toInteger(row.DayofMonth)}]-(destination)
ON CREATE SET connection.departure = toInteger(row.CRSDepTime),
            connection.arrival = toInteger(row.CRSArrTime)
```

## Exercise 2: Loading Airport Data

In the query edit pane of Neo4j Browser, execute the browser command:

```
:play 4.0-neo4j-modeling-exercises
```

and follow the instructions for Exercise 2.

Note: This exercise has 9 steps. Estimated time to complete: 30 minutes.



# Summary

You should now be able to:

- Write Cypher code to import CSV data with Cypher
- Confirm that the data has been loaded



# Profiling Queries

# In This Module You'll Learn ...

At the end of this module, you should be able to:

- Profile queries against the graph
- Determine if a query can be improved



# Profiling a Query

PROFILE

MATCH (origin:Airport)-

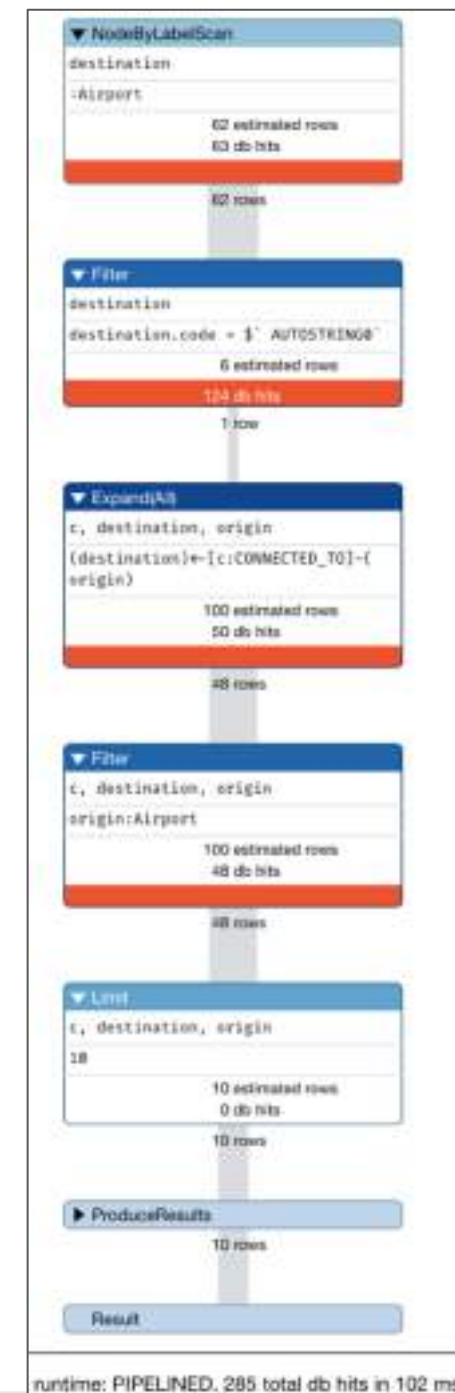
[c:CONNECTED\_TO]->(destination:Airport)

WHERE destination.code = 'LAS'

RETURN origin, destination, c LIMIT 10

Code to profile a query

- Retrieve all connections that have a destination of *LAS*




# Analyzing the Query Profile (1 of 3)

## PROFILE

```
MATCH (origin:Airport)-  
      [c:CONNECTED_TO]->(destination:Airport)  
WHERE destination.code = 'LAS'  
RETURN origin, destination, c LIMIT 10
```

## Code analysis

- destination:Airport
  - 62 total airports
- Destination.code = 'LAS'
  - One airport with code LAS



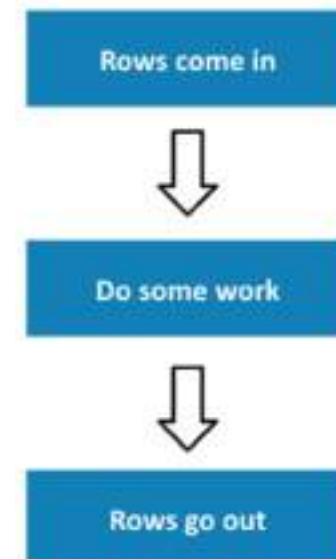
# Analyzing the Query Profile (2 of 3)

## PROFILE

```
MATCH (origin:Airport)-  
      [c:CONNECTED_TO]->(destination:Airport)  
WHERE destination.code = 'LAS'  
RETURN origin, destination, c LIMIT 10
```

## Code analysis

- Pattern `origin-[c]->(destination)`
  - 48 nodes
- No filter for pattern
  - All 48 nodes pass through



# Analyzing the Query Profile (3 of 3)

## PROFILE


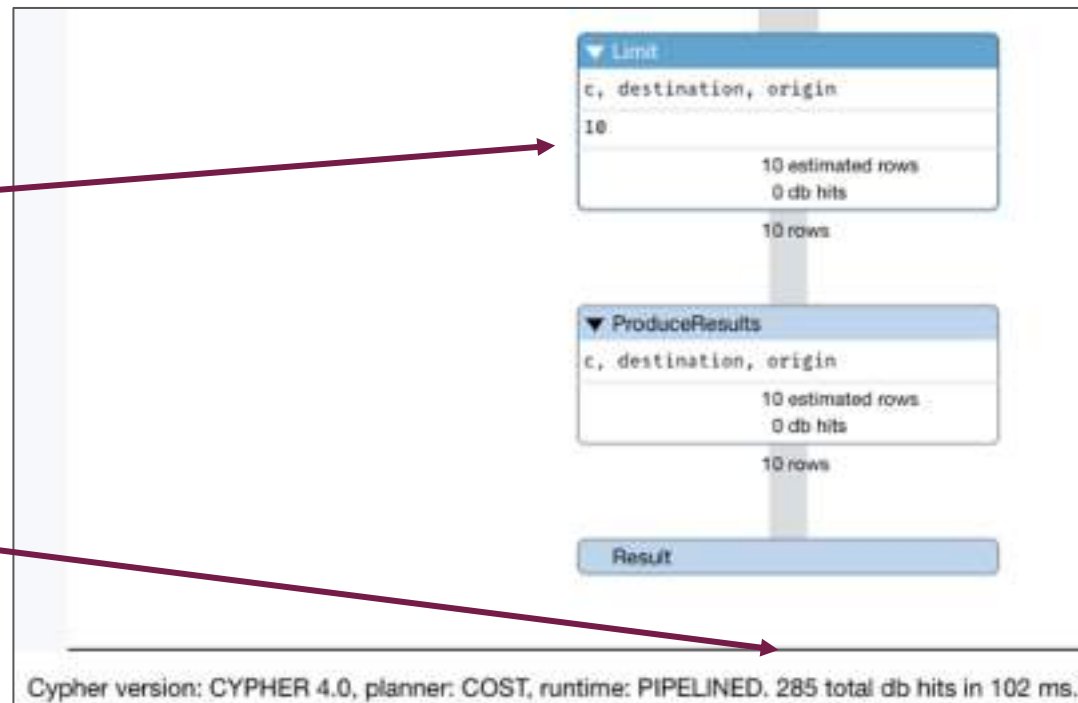
```
MATCH (origin:Airport)-  
      [c:CONNECTED_TO]->(destination:Airport)  
WHERE destination.code = 'LAS'  
RETURN origin, destination, c LIMIT 10
```

## Code analysis

- Product results
  - 48 nodes
- Total cost
  - 285 db hits

## Code analysis

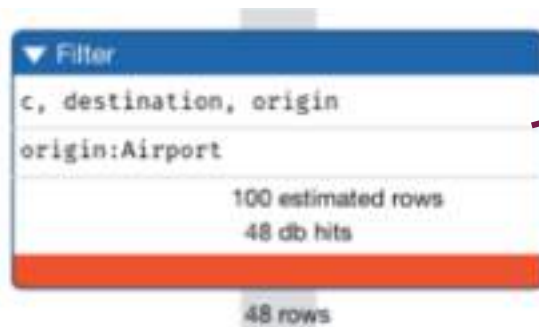
- LIMIT 10
- Product results
  - 48 nodes
- Total cost
  - 285 db hits



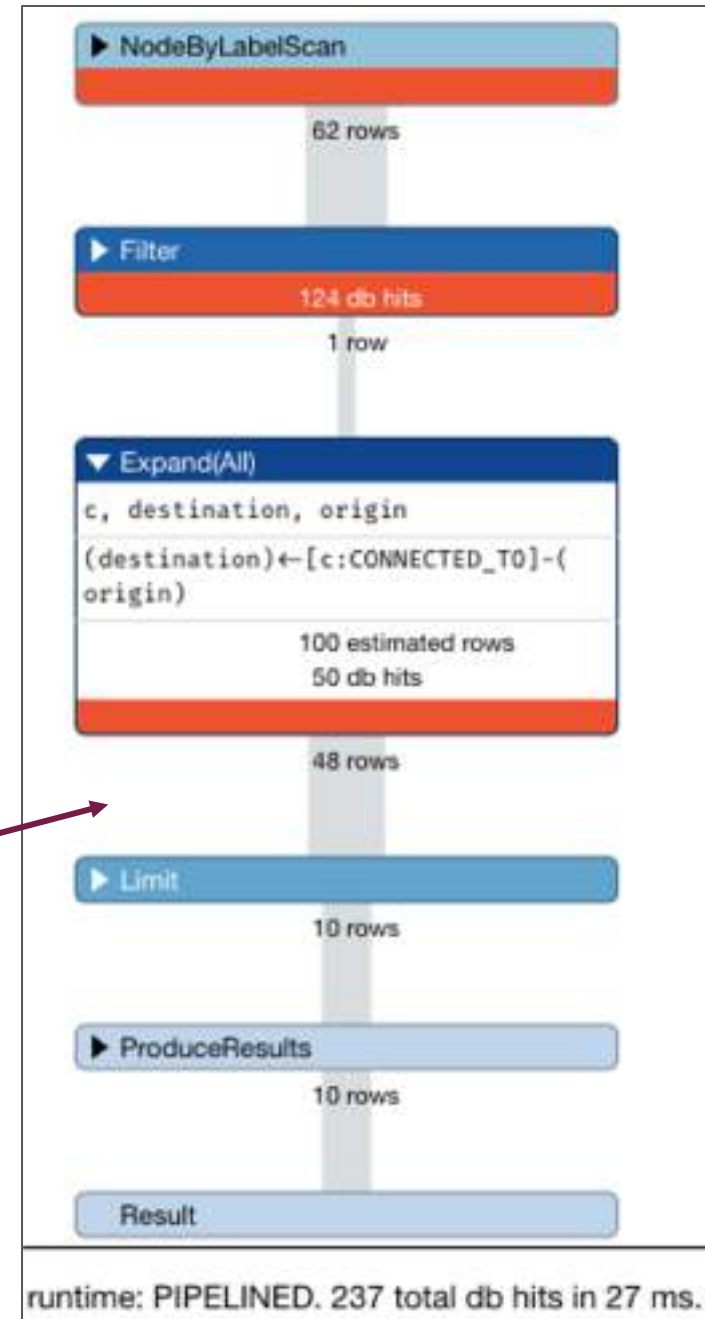
# Analyzing Without :Airport Filter

```
PROFILE  
MATCH (origin)-  
      [c:CONNECTED_TO]->(destination:Airport)  
WHERE destination.code = 'LAS'  
RETURN origin, destination, c LIMIT 10
```

← :Airport label  
not specified



Without this Filter  
step there are 48  
less db hits





# Exercise 3: Profiling Queries

In the query edit pane of Neo4j Browser, execute the browser command:

```
:play 4.0-neo4j-modeling-exercises
```

and follow the instructions for Exercise 3.

Note: This exercise has 2 steps. Estimated time to complete: 15 minutes.



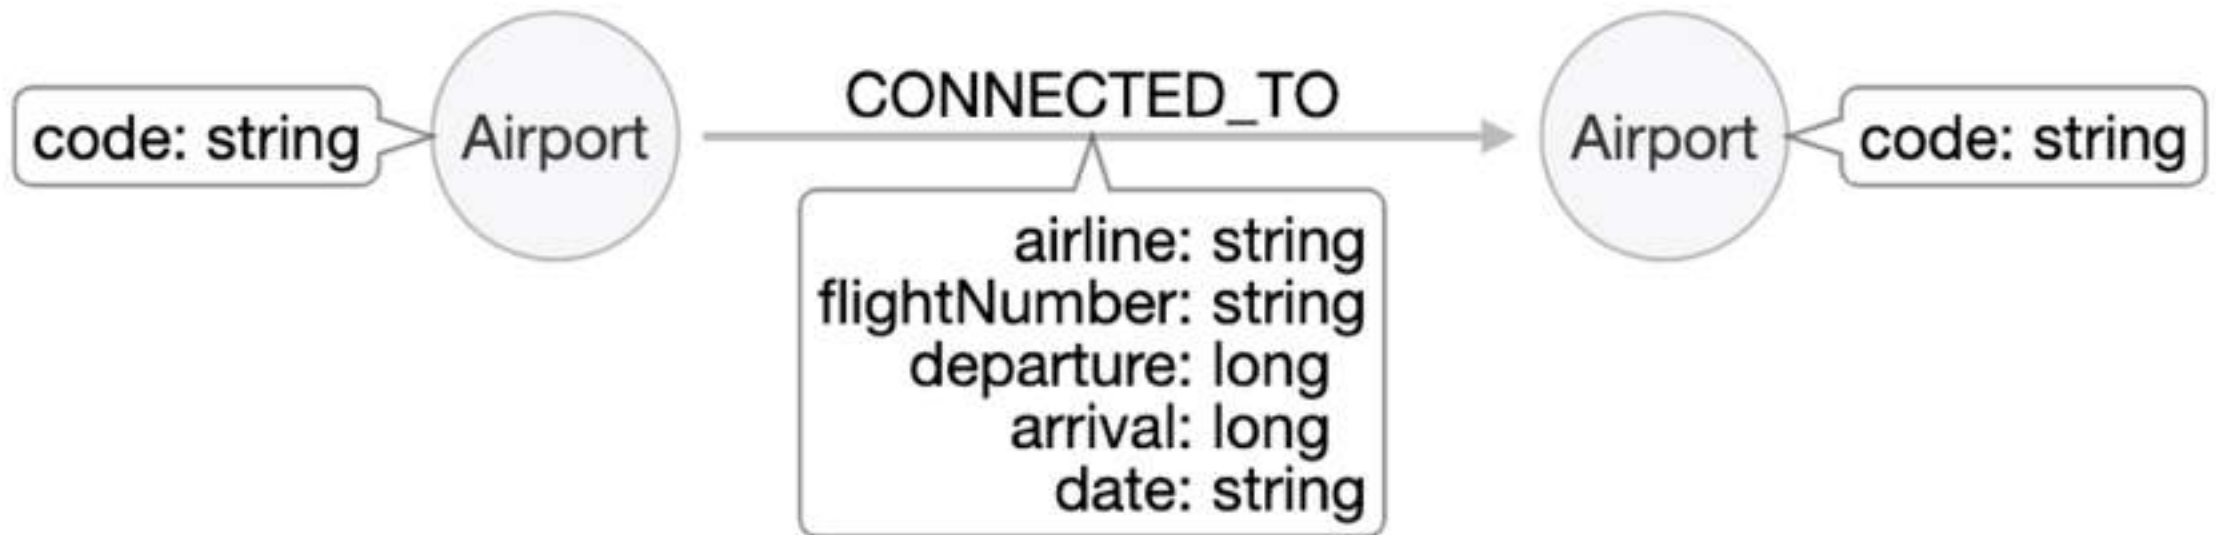


# Does the Model Need to be Changed?

The previous exercise asked this question:

- What are the airports and flight information for flight number 1016 for airline WN?

The current model:



# Analyzing the Other Question

What are the airports and flight information for flight number 1016 for airline WN?

```
PROFILE
MATCH (origin:Airport)-
      [connection:CONNECTED_TO]->
      (destination:Airport)
WHERE connection.airline = 'WN'
      AND connection.flightNumber = '1016'
RETURN origin.code, destination.code,
       connection.date,
       Connection.departure,
       connection.arrival
```



# Summary

You should now be able to:

- Profile queries against the graph
- Determine if a query can be improved

# Refactoring Graphs

# Refactoring Steps

1. Create constraints as needed
2. Execute the refactor:
  - a. MATCH the data to be moved
  - b. Create new nodes
  - c. Create new relationships
3. Create indexes as needed
4. PROFILE all queries against the new model

If the **new model** performs well for **all** queries

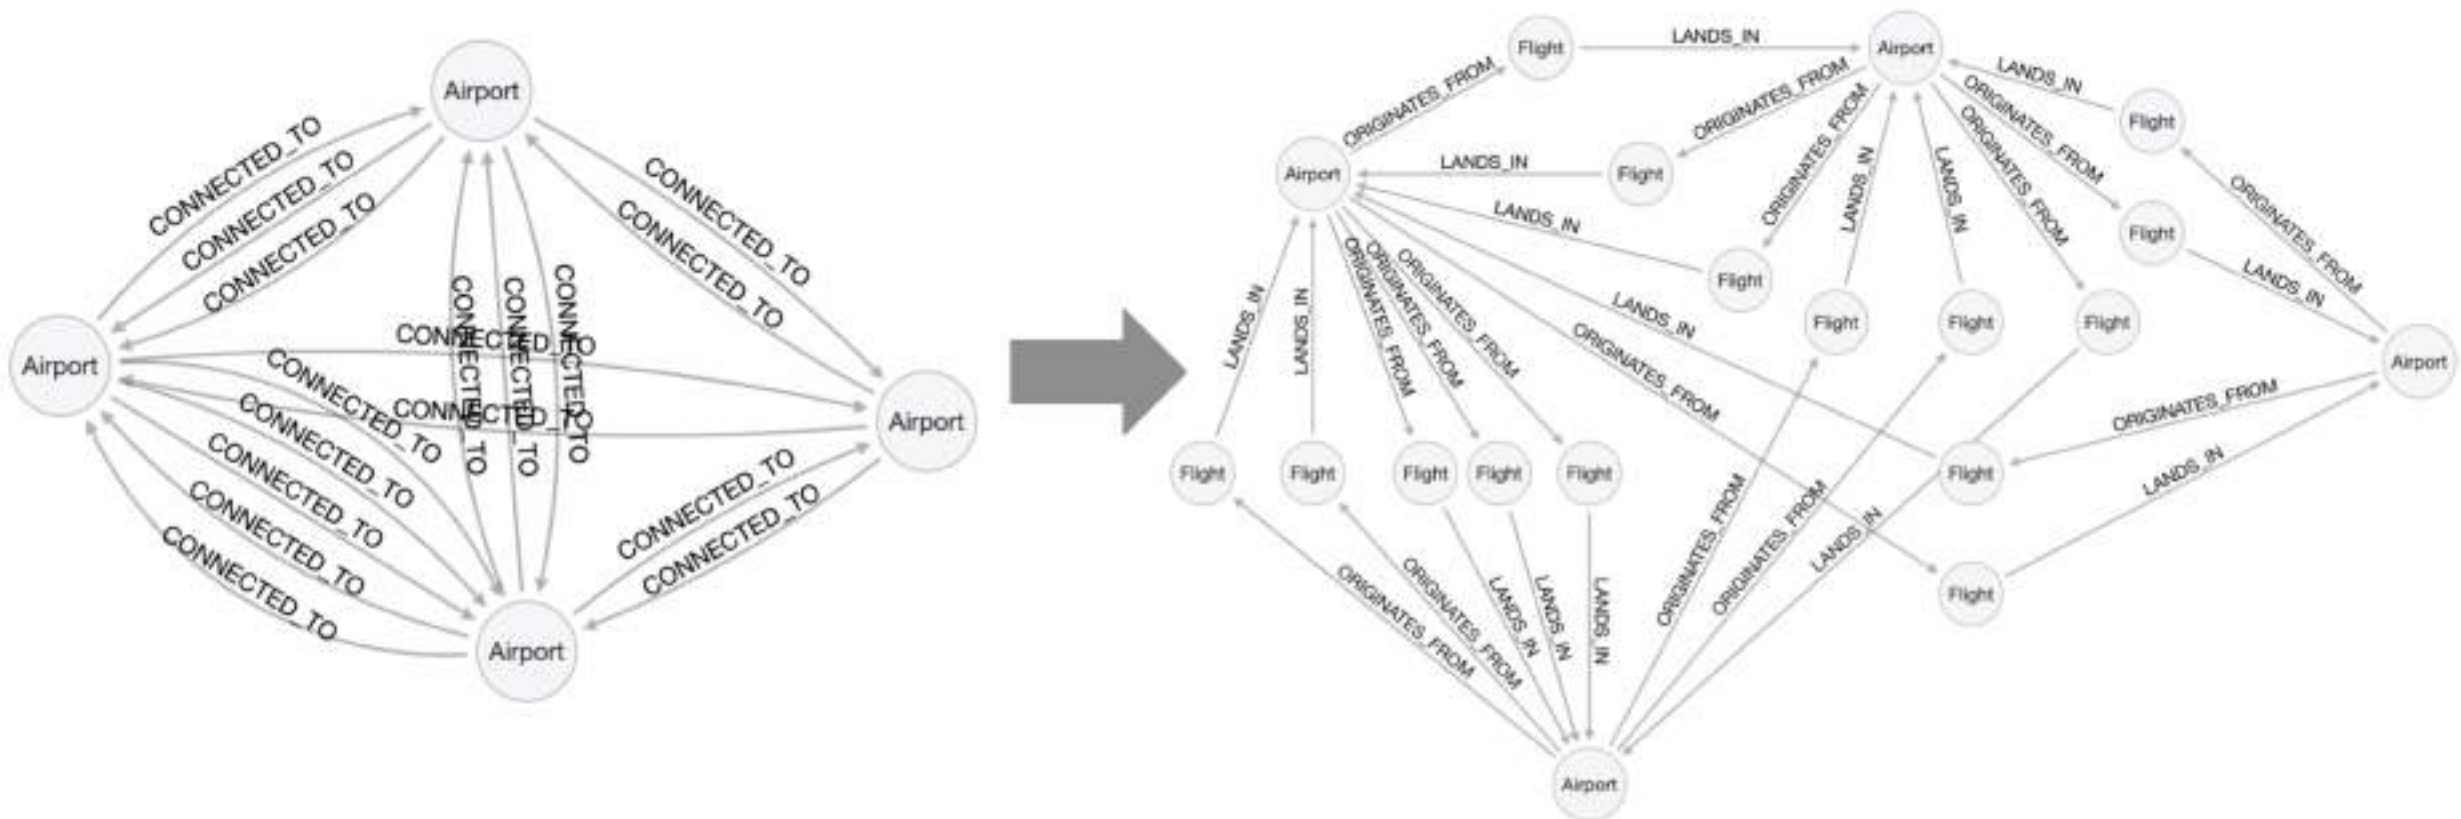
- **Delete** the **old model**
- Otherwise,  
keep **both models**



# Evolving the Model

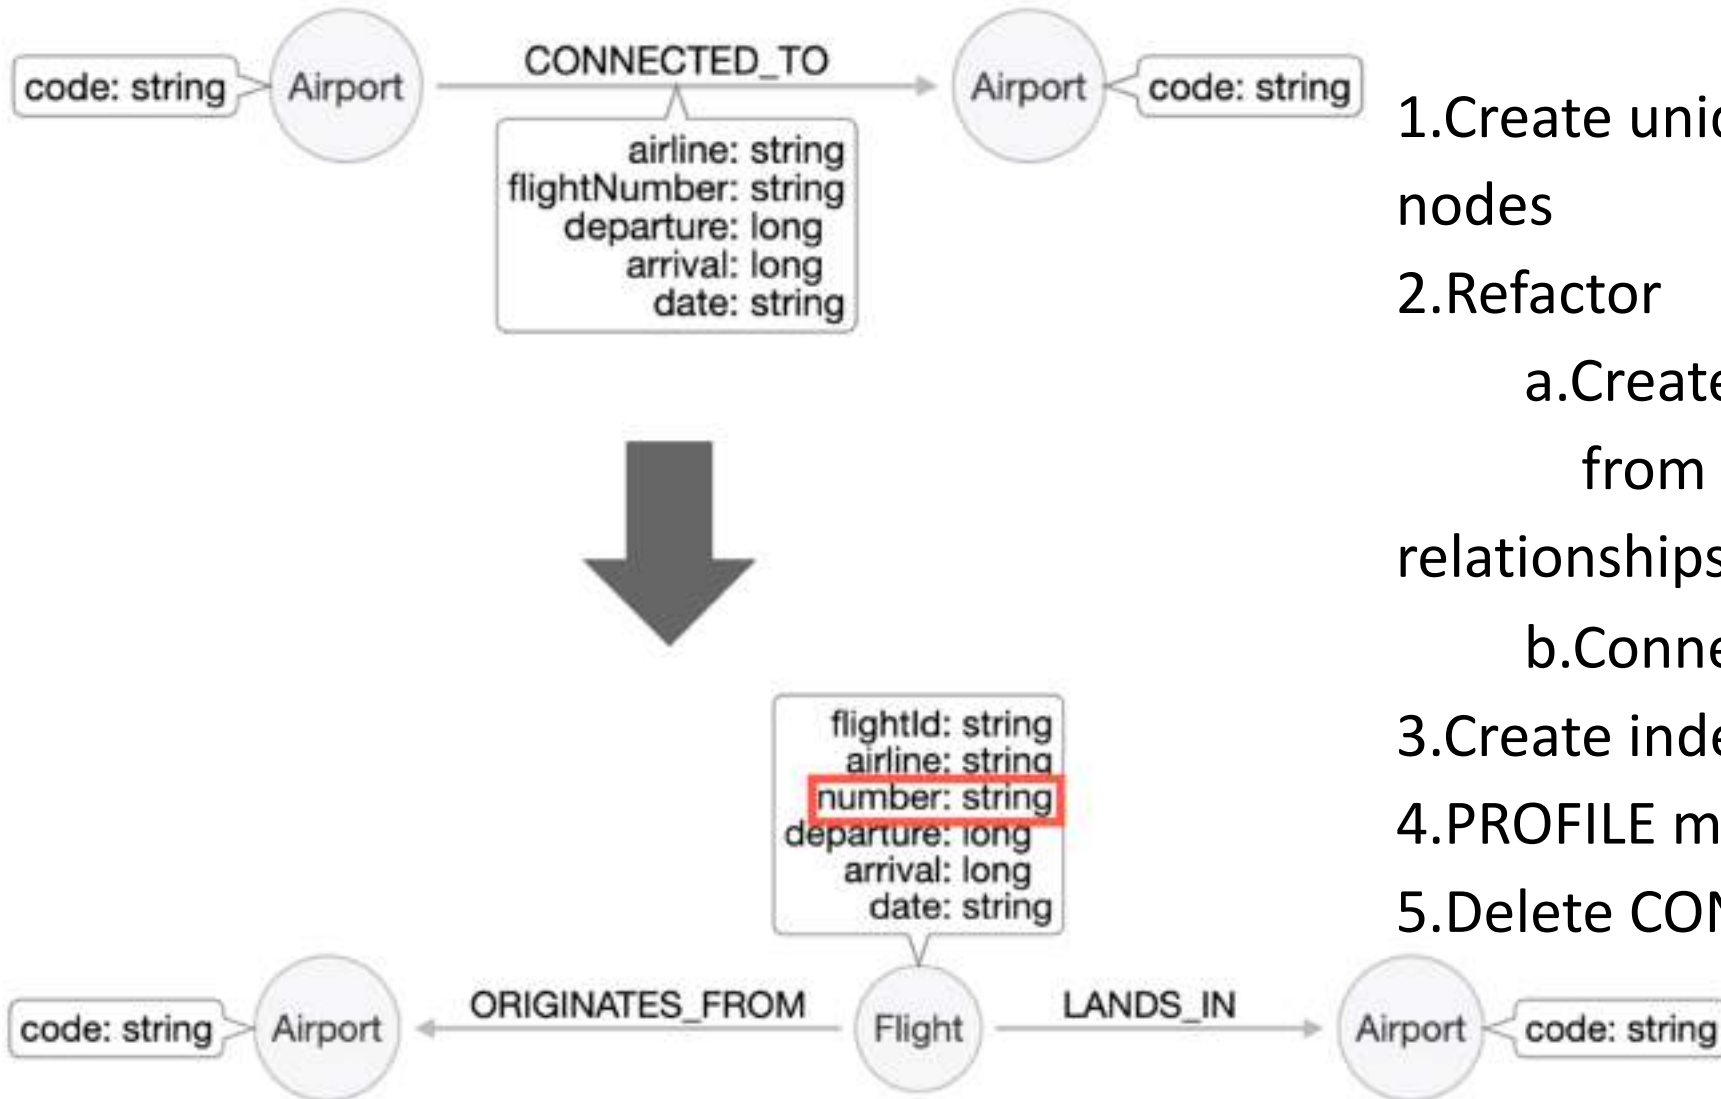
To solve this problem intermediate node Flight is introduced

- It is based upon the properties of the *CONNECTED\_TO* relationship





# Refactor Details



1.Create unique constraint on Flight nodes

2.Refactor

a.Create Flight nodes  
from CONNECTED\_TO  
relationships

b.Connect Flights to Airports

3.Create index on Flight.number

4.PROFILE modified queries

5.Delete CONNECTED\_TO relationships

# Unique flightId

The *Flight.flightId* property is composed of five pieces of data

- This combination assures that all *Flight* nodes are **unique**
  1. Airline
  2. flightNumber
  3. Code for the origin Airport
  4. Code for the destination Airport
  5. Date

# Creating a Constraint

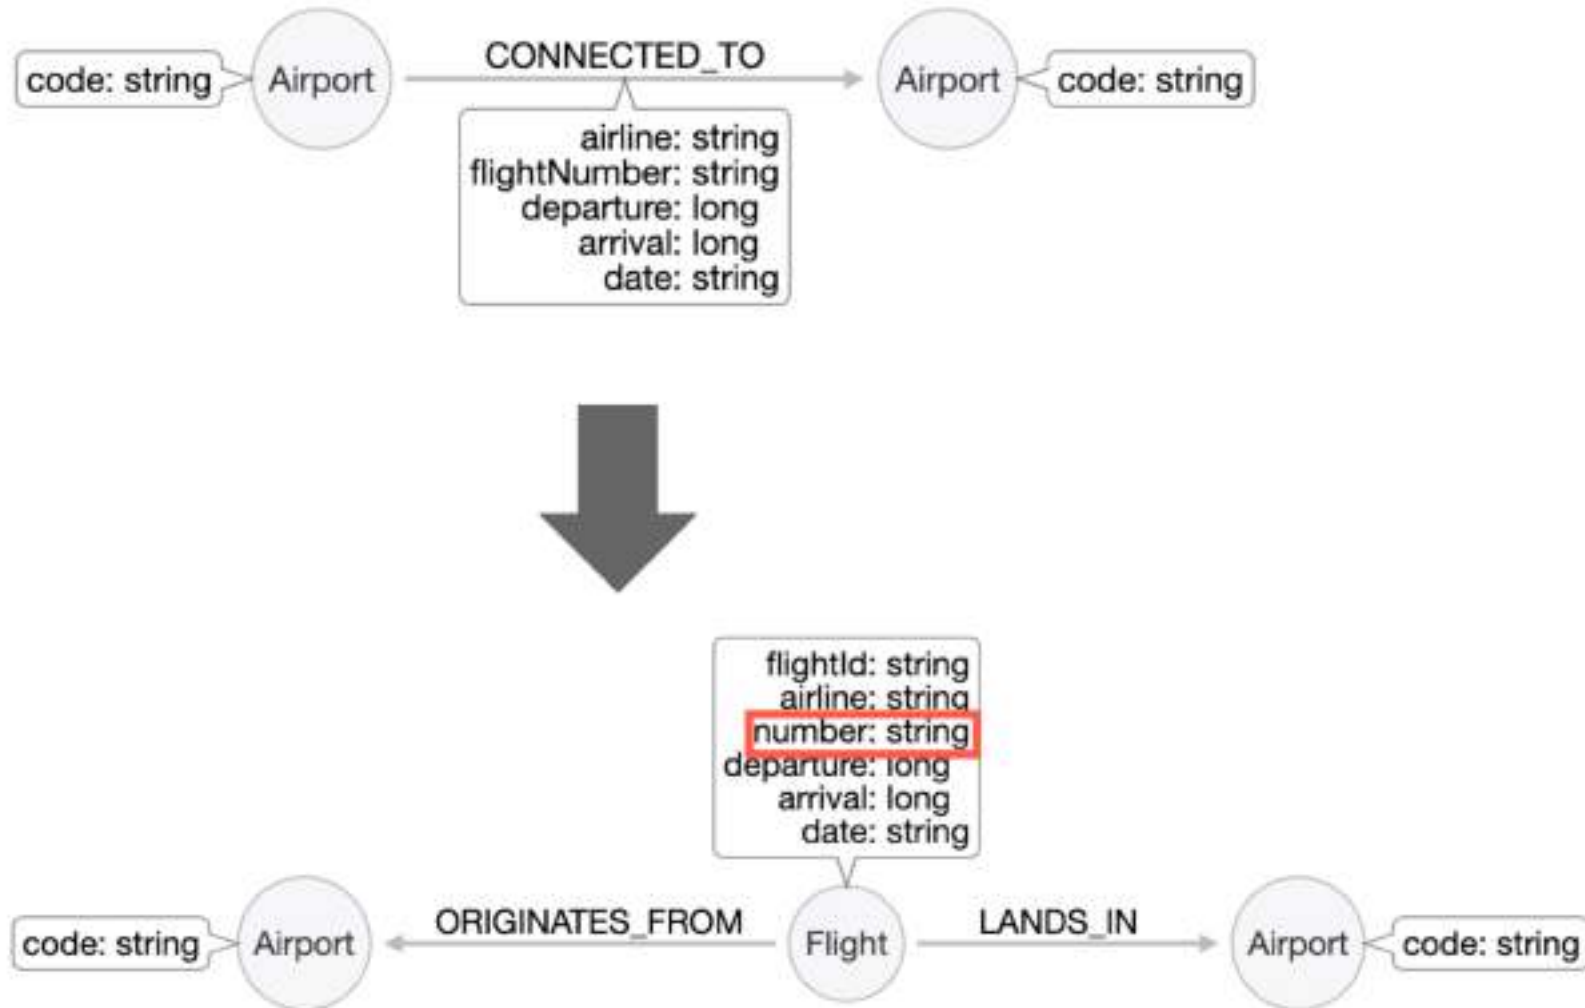
Prior to refactoring the uniqueness constraint is added

```
CREATE CONSTRAINT Flight_flightId_constraint ON (f:Flight)  
  ASSERT f.flightId IS UNIQUE
```

# Refactoring the Graph

```
MATCH (origin:Airport)-[connection:CONNECTED_TO]-
>(destination:Airport)
MERGE (newFlight:Flight {flightId: connection.airline +
    connection.flightNumber +
        '_' + connection.date + '_' + origin.code + '_' +
    destination.code })
ON CREATE SET newFlight.date = connection.date,
    newFlight.airline = connection.airline,
    newFlight.number = connection.flightNumber,
    newFlight.departure = connection.departure,
    newFlight.arrival = connection.arrival
MERGE (origin)<-[:ORIGINATES_FROM]-(newFlight)
MERGE (newFlight)-[:LANDS_IN]->(destination)
```

# Two Models Exist In The Same Graph



# Create an Index on Flight Number

```
CREATE INDEX Flight_number_index FOR (f:Flight) ON (f.number)
```



# Profiling the Query

Question: What are the airports and flight information for flight number 1016 for airline WN?

Original  
Query

```
PROFILE
MATCH (origin:Airport)-[connection:CONNECTED_TO]-
>(destination:Airport)
WHERE connection.airline = 'WN' AND
connection.flightNumber = '1016'
```

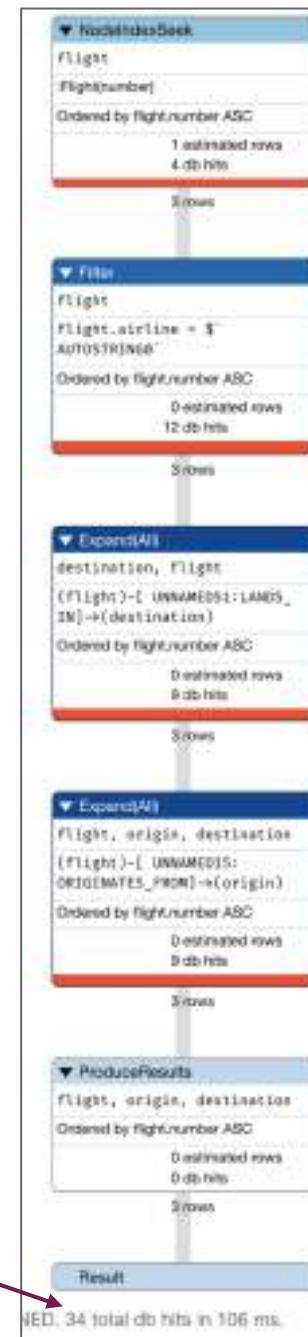
New Query

```
PROFILE
MATCH (origin)<-[:ORIGINATES_FROM]-(flight:Flight)-
[:LANDS_IN]->(destination)
WHERE flight.airline = 'WN' AND
flight.number = '1016' RETURN origin, destination,
flight
```

# Profile Result



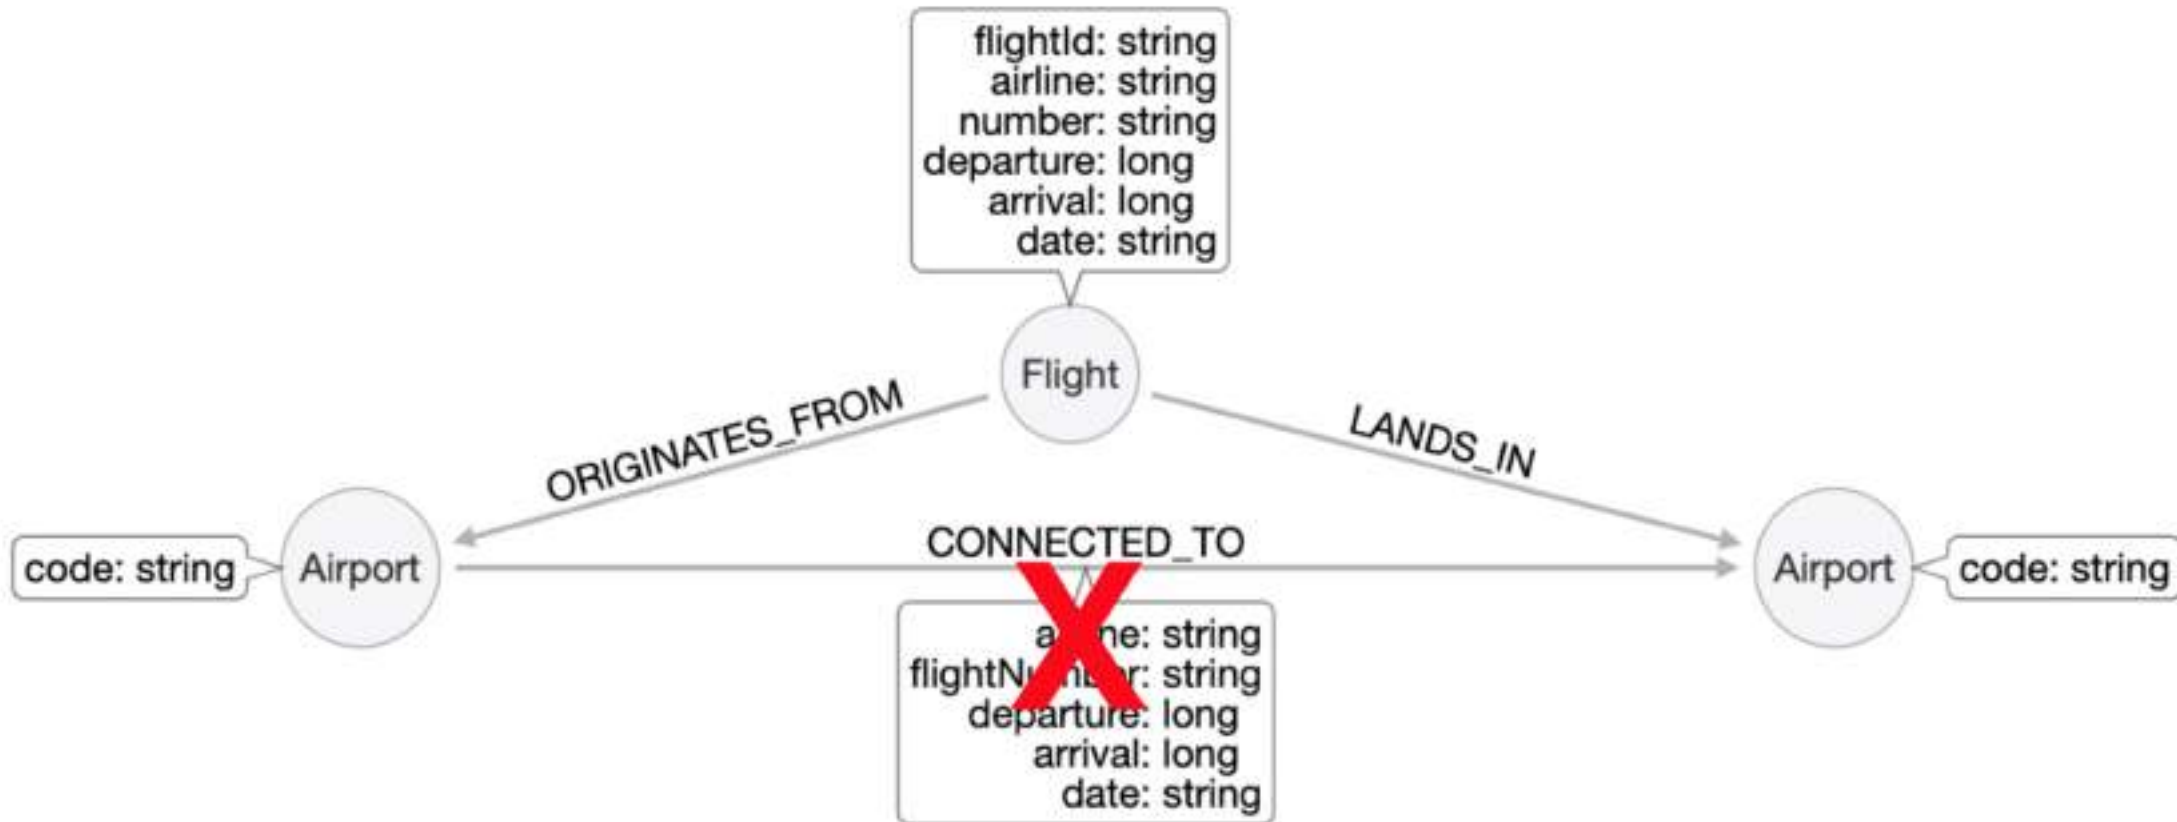
D. 34 total db hits in 106 ms.



4ED. 34 total db hits in 106 ms.

# Removing the Old Model

```
MATCH ()-[connection:CONNECTED_TO]->()  
DELETE connection
```



## Exercise 4: Creating Flight Nodes from CONNECTED\_TO Relationships

In the query edit pane of Neo4j Browser, execute the browser command:

```
:play 4.0-neo4j-modeling-exercises
```

and follow the instructions for Exercise 4.

Note: This exercise has 7 steps. Estimated time to complete: 30 minutes.



# An Additional Domain Question

Here is another question is that needs to be answered by the application:

- As a frequent traveller I want to find **flights** from **<origin>** to **<destination>** on **<date>** so that I can book my business flight

For example:

- Find all **flights** going from Los Angeles (**LAS**)  
to Chicago Midway International (**MDW**)  
on **January 3rd, 2019**.

# Implementing the Query

```
MATCH (origin:Airport {code: 'LAS'})  
      <-[:ORIGINATES_FROM]-(flight:Flight)-[:LANDS_IN]->  
      (destination:Airport {code: 'MDW'})  
WHERE flight.date = '2019-1-3'  
RETURN origin, destination, flight
```

The next exercise makes use of this query

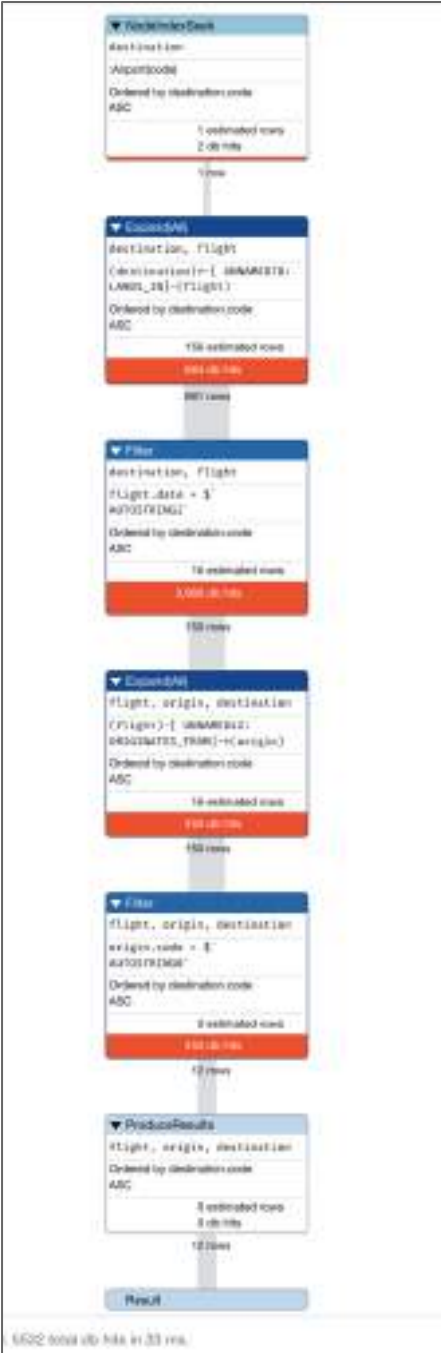
- From the US Bureau of Transportation data 10k nodes will be added to the graph



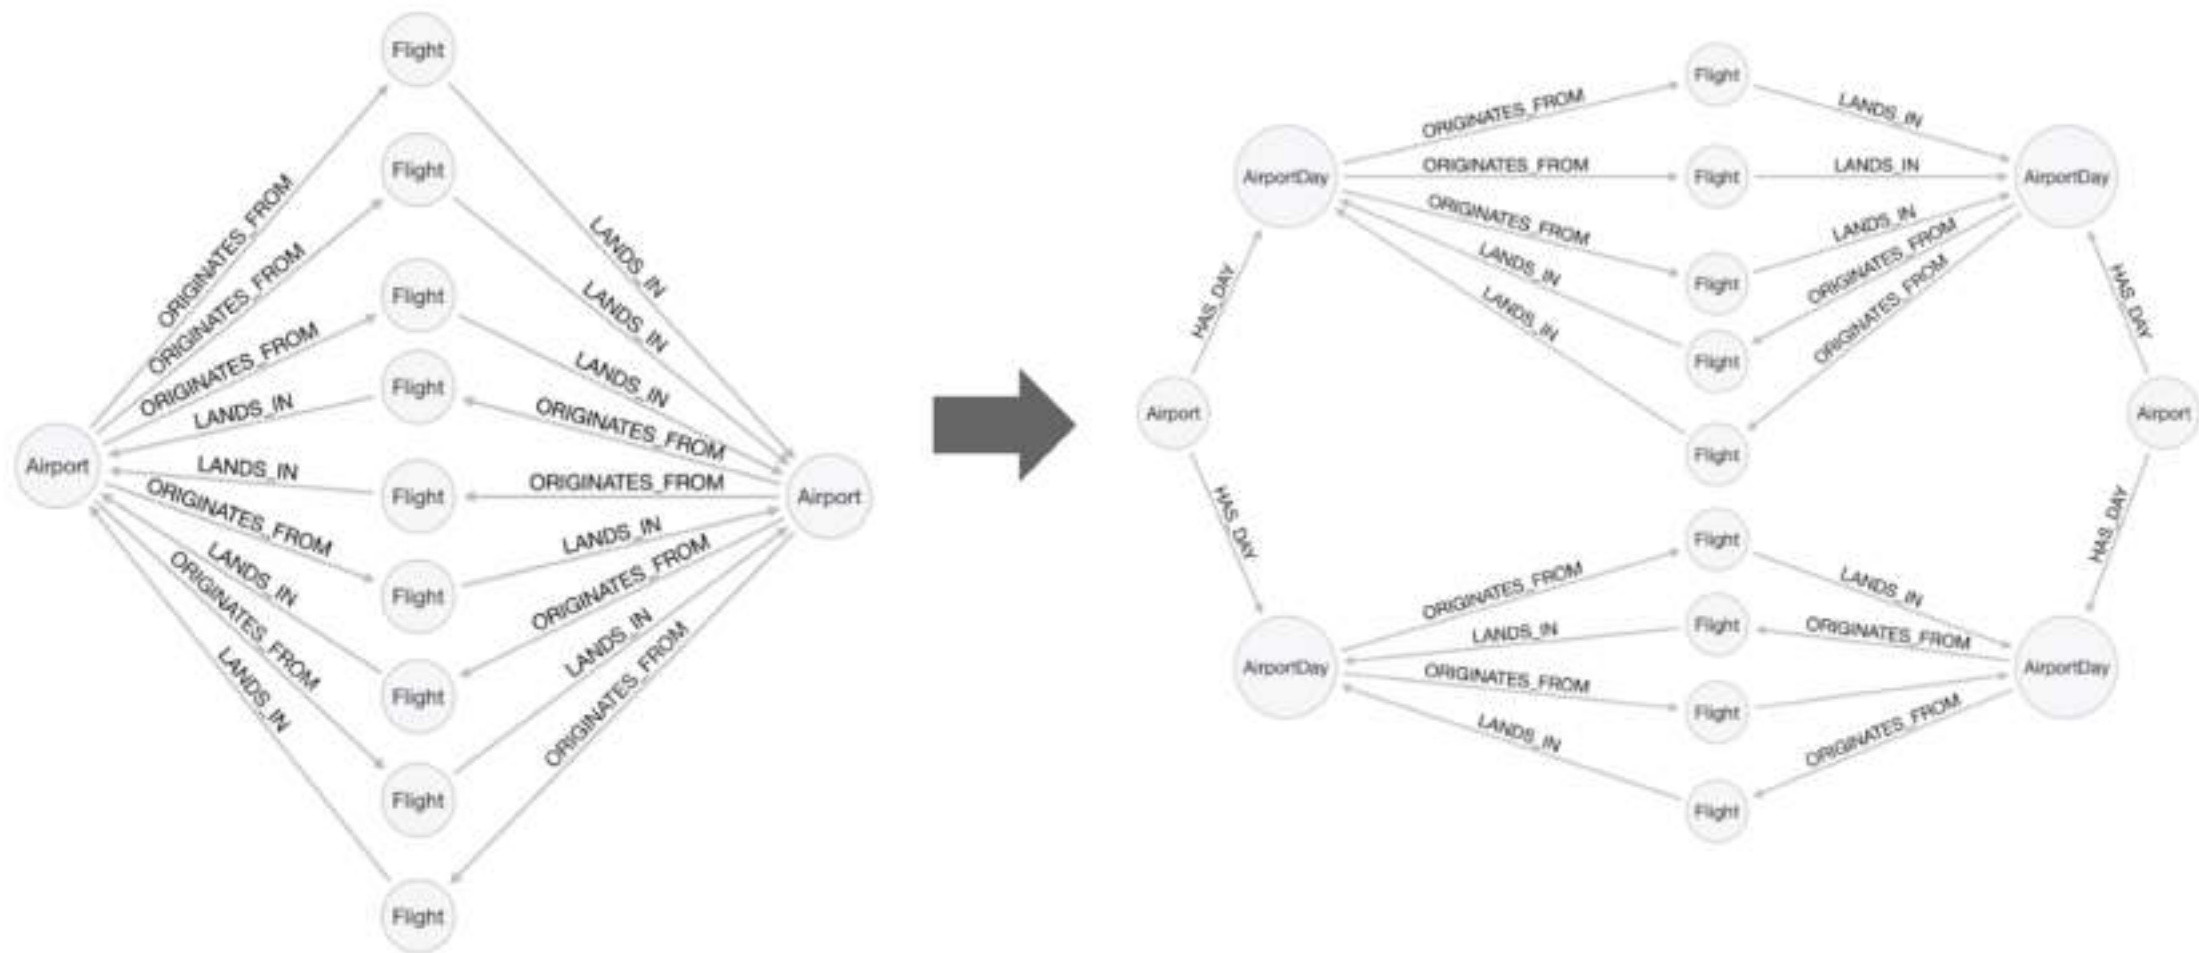
# Profiling the Query



The result after profiling the query that contains 10k nodes



# Performing Another Refactor



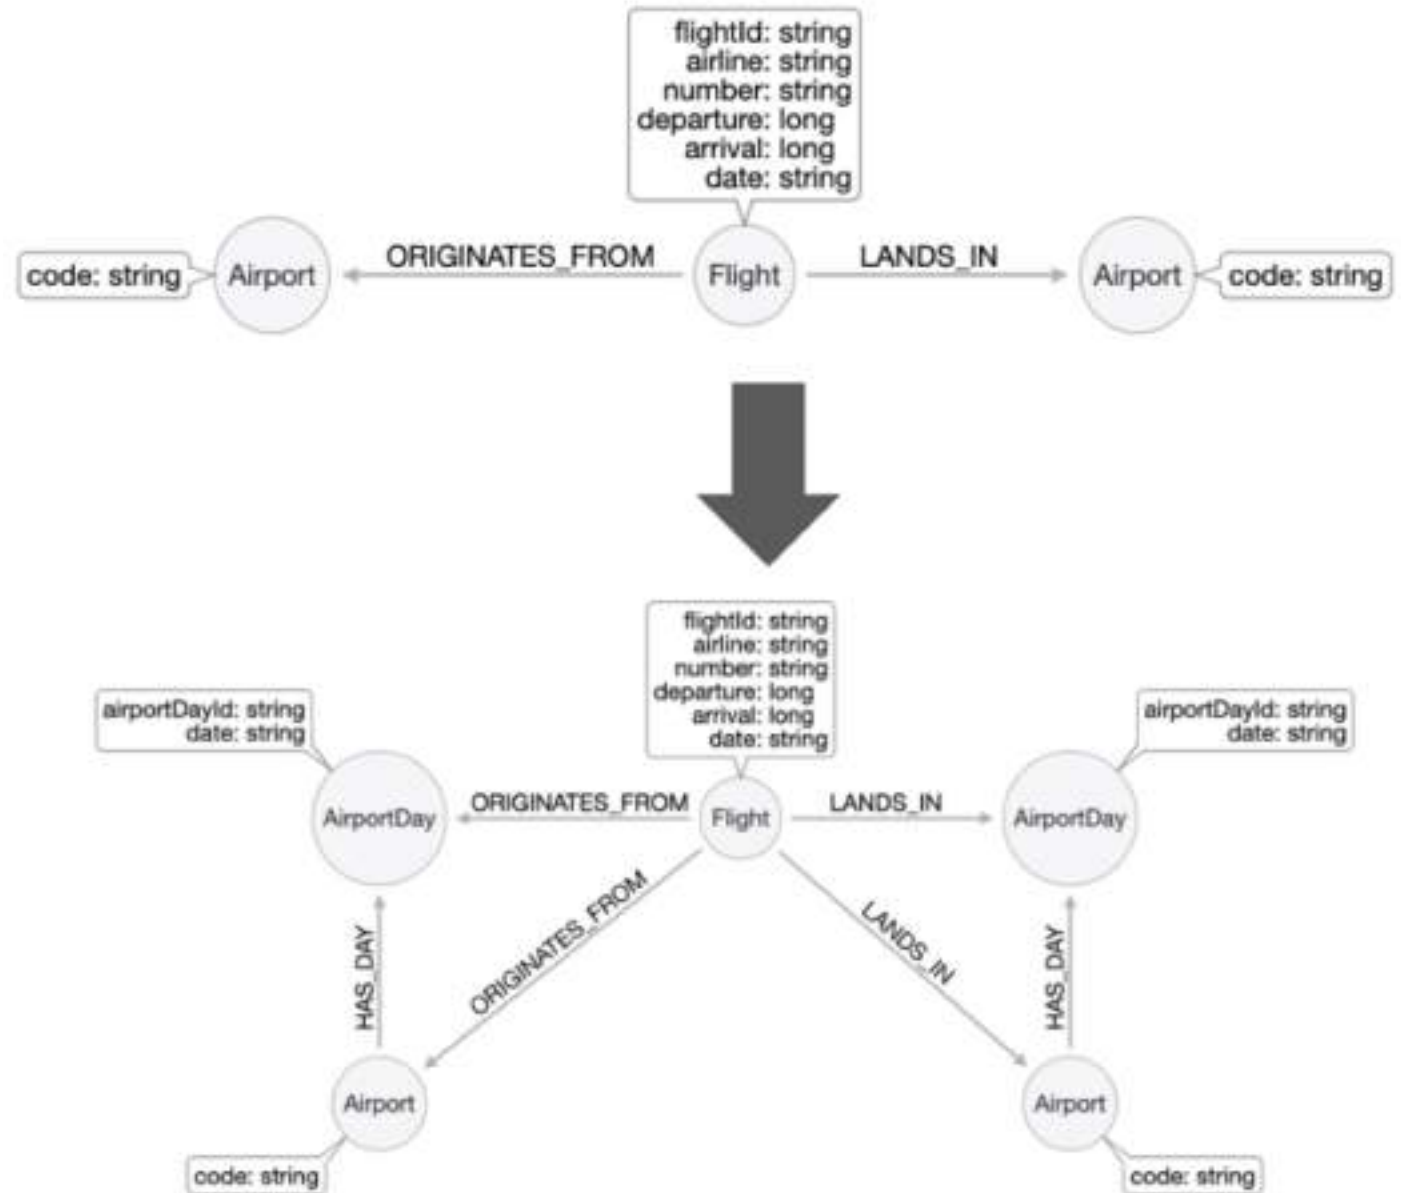
Note: A node is only pulled out if you are going to query through it, otherwise a property will suffice

# Refactor Details

## CREATE CONSTRAINT

AirportDay\_airportDayId\_constraint

```
ON (a:AirportDay)
ASSERT a.airportDayId IS
UNIQUE
```



# Refactor Implementation

```
MATCH (origin:Airport)<-[:ORIGINATES_FROM]-(flight:Flight)-
      [:LANDS_IN]->(destination:Airport)
MERGE (originAirportDay:AirportDay {airportDayId: origin.code +
  '_' + flight.date})
SET originAirportDay.date = flight.date
MERGE (destinationAirportDay:AirportDay
      {airportDayId: destination.code + '_' + flight.date})
SET destinationAirportDay.date = flight.date
MERGE (origin)-[:HAS_DAY]->(originAirportDay)
MERGE (flight)-[:ORIGINATES_FROM]->(originAirportDay)
MERGE (flight)-[:LANDS_IN]->(destinationAirportDay)
MERGE (destination)-[:HAS_DAY]->(destinationAirportDay)
```

# Profiling the First Query



PROFILE

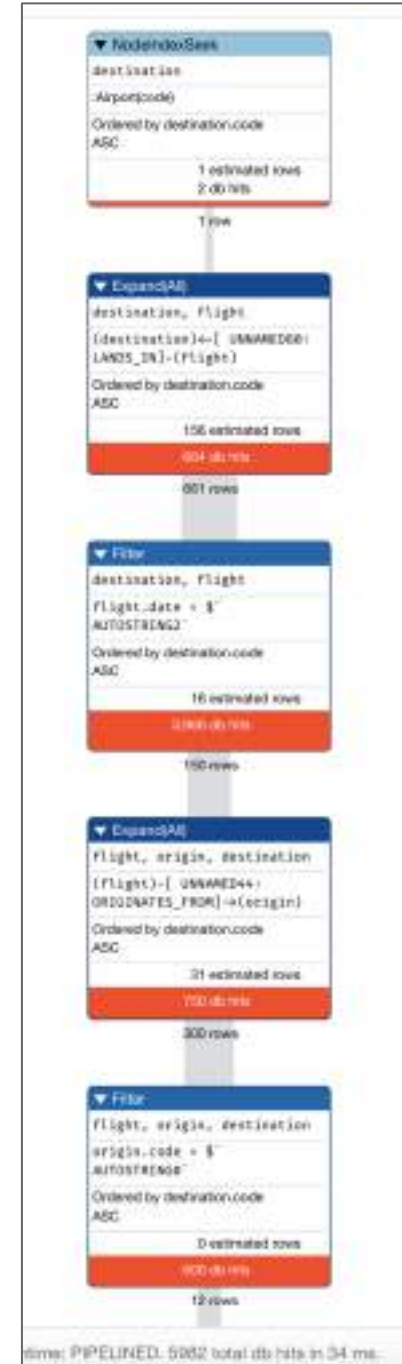
```
MATCH (origin)-[:ORIGINATES_FROM]-  
(flight:Flight)-
```

```
[:LANDS_IN]->(destination)
```

```
WHERE flight.airline = 'WN' AND  
flight.number = '1016' RETURN origin,  
destination, flight
```

# Profiling the Original Second Query

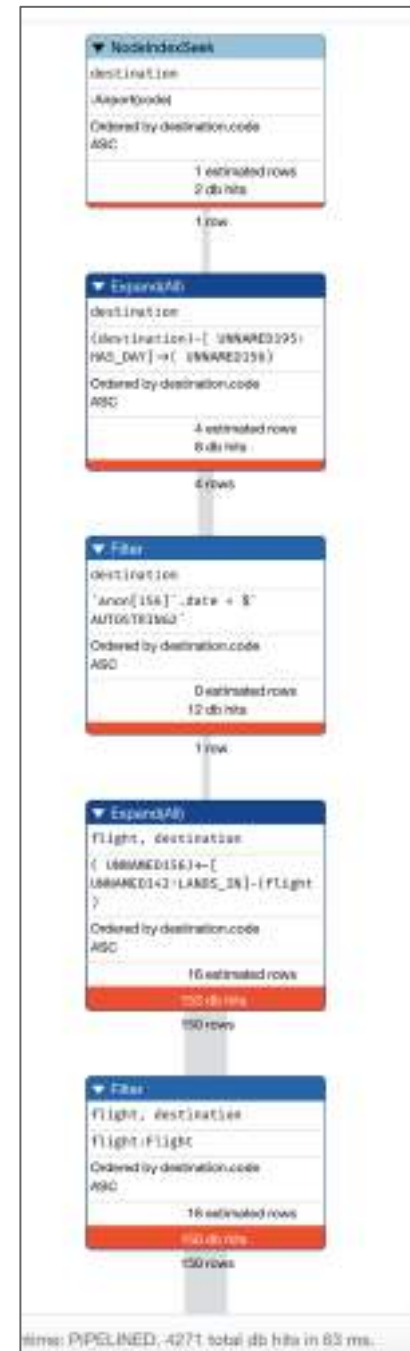
```
PROFILE MATCH (origin:Airport
{code: 'LAS'})
  <-[:ORIGINATES_FROM]-
(flight:Flight)-
[:LANDS_IN]->
(destination:Airport {code:
'MDW'})
WHERE flight.date = '2019-1-3'
RETURN origin, destination, flight
```





# Profiling the Revised Second Query

```
PROFILE MATCH (origin:Airport {code:
'LAS'})-
[:HAS_DAY]->(:AirportDay {date:
'2019-1-3'})<-
[:ORIGINATES_FROM]-(flight:Flight),
(flight)-[:LANDS_IN]->
(:AirportDay {date: '2019-1-3'})<-
[:HAS_DAY]-(destination:Airport
{code: 'MDW'})
RETURN origin, destination, flight
```



## Exercise 5: Creating the AirportDay Node From the Airport and Flight Nodes

In the query edit pane of Neo4j Browser, execute the browser command:

```
:play 4.0-neo4j-modeling-exercises
```

and follow the instructions for Exercise 5.

Note: This exercise has 7 steps. Estimated time to complete: 30 minutes.



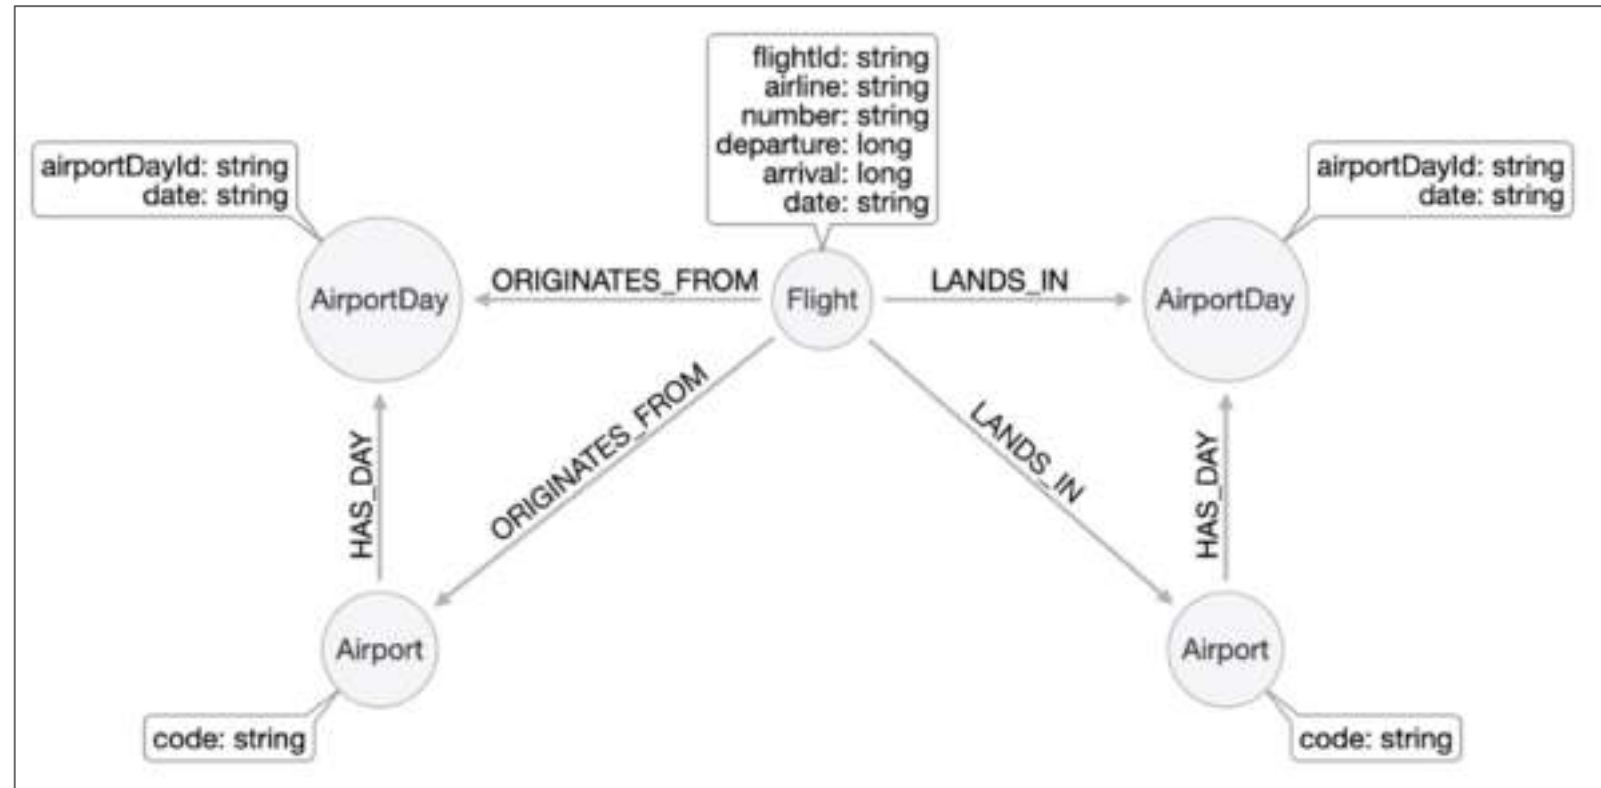
# More Questions for the Model?

The model performs well  
for these questions:

1. What are the airports and flight information for flight number xx for airline yy?
2. Find all the flights going from xx to yy on the date zz.

What if this question is added:

- Which airport has the most incoming flights?



# Another Question for the Model

Suppose we added this question:

- What are the flights from LAS that arrive at MDW on 2019-1-3?

## PROFILE

```
MATCH (origin:Airport {code: 'LAS'})-[:HAS_DAY]->
      (originDay:AirportDay),
      (originDay)<-[:ORIGINATES_FROM]-(flight:Flight),
      (flight)-[:LANDS_IN]->(destinationDay),
      (destinationDay:AirportDay)<-[:HAS_DAY]-
      (destination:Airport {code: 'MDW'})
WHERE originDay.date = '2019-1-3' AND
      destinationDay.date = '2019-1-3'
RETURN flight.date, flight.number, flight.airline,
       flight.departure, flight.arrival
ORDER BY flight.date, flight.departure
```

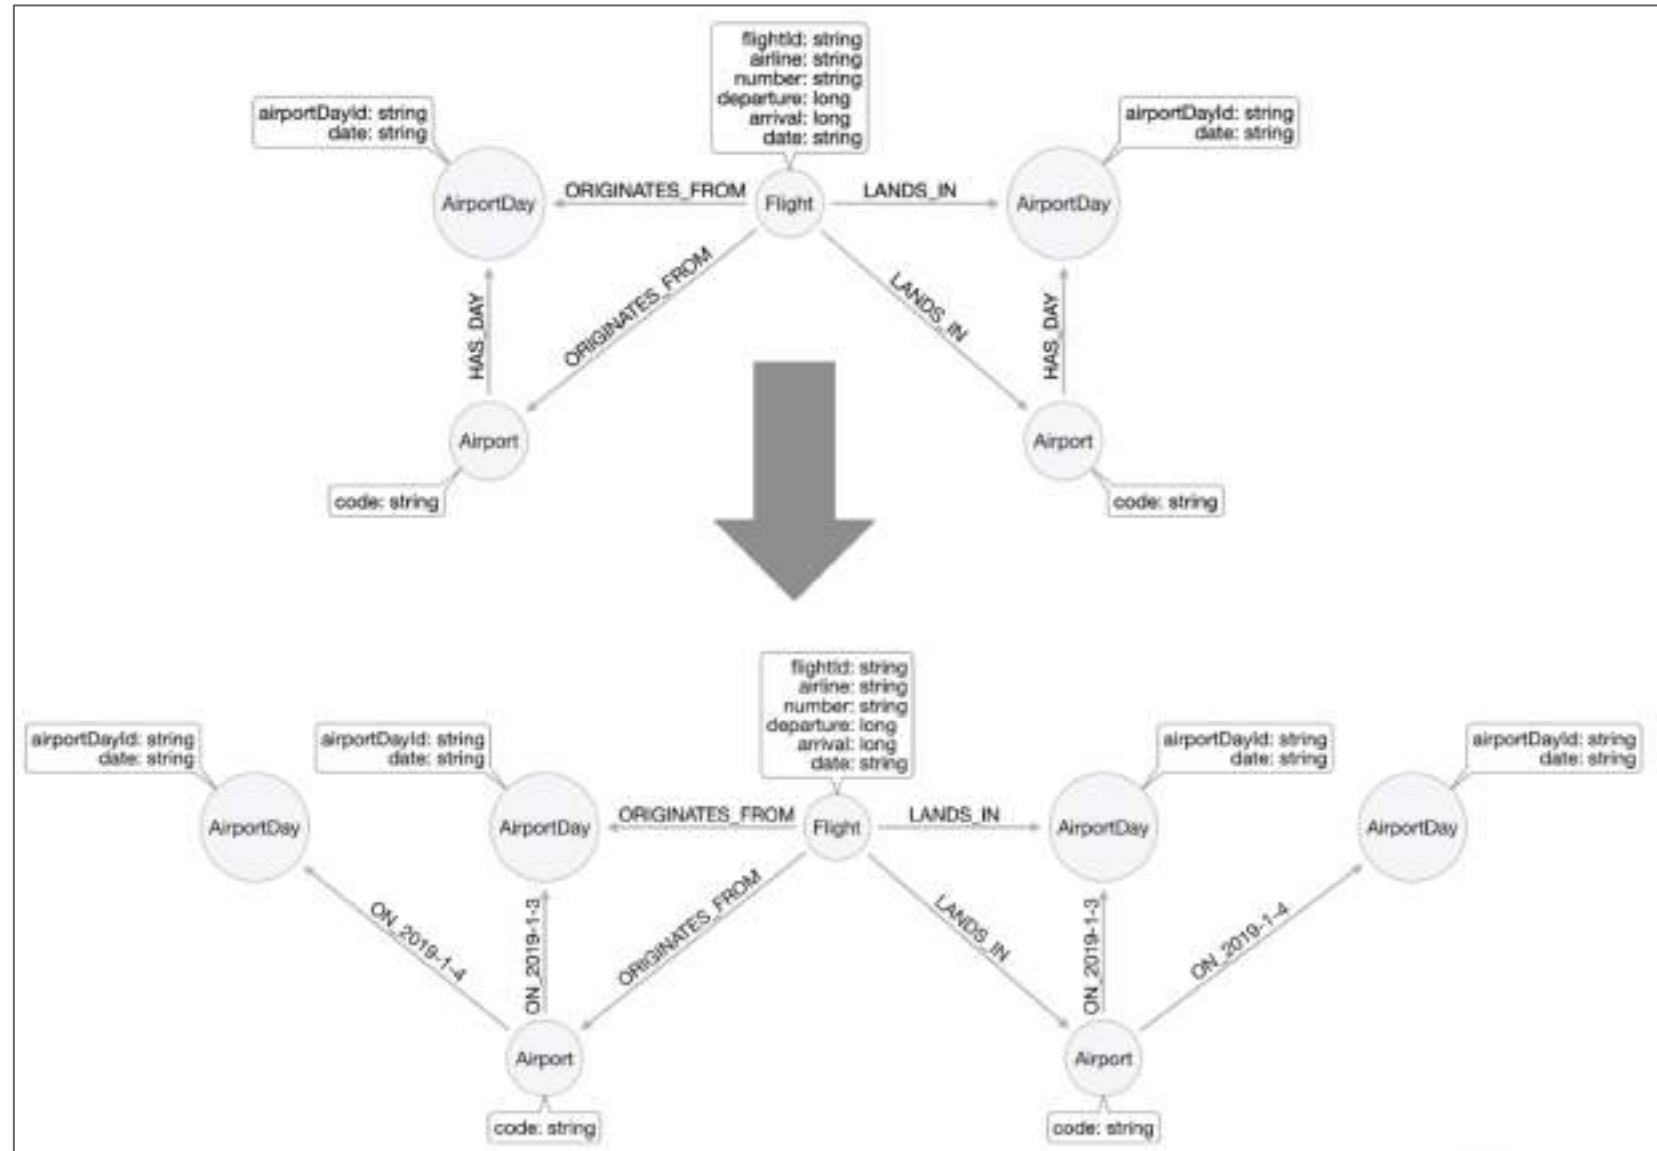


# Refactoring for Specific Relationships

A best practice for graph data modeling:

- Make relationships more specific if it helps query performance

Modify *HAS\_DAY* relationship to be *ON\_2019-1-3*, *ON\_2019-1-4*, etc.



# APOC to the Rescue!

APOC supports creating relationships based upon data in the graph

syntax:

```
apoc.create.relationship(startNode(<relationship-variable>),  
                        '<new-relationship-value>',  
                        {<relationship-property list>},  
                        endNode(<relationship-variable>)  
                        )  
YIELD rel
```



# Creating Specialized Relationships with APOC

Code to transform *HAS\_DAY* relationships to specific relationships:

```
MATCH (origin:Airport)-[hasDay:HAS_DAY]->(ad:AirportDay)
CALL apoc.create.relationship(startNode(hasDay),
                             'ON_' + ad.date,
                             {},
                             endNode(hasDay) ) YIELD rel
RETURN COUNT(*)
```

# Refactoring Result

Use database

neo4j - default

Node Labels

\*(10,319) Airport AirportDay Flight

Relationship Types

\*(40,510) HAS\_DAY LANDS\_IN ON\_2019-1-3 ON\_2019-1-4 ON\_2019-1-5 ON\_2019-1-6 ORIGINATES\_FROM

```
neo4j$ MATCH (origin:Airport)-[hasDay:HAS_DAY]→(ad:AirportDay) CALL apoc.create.rela...
```

	COUNT(*)
A	255

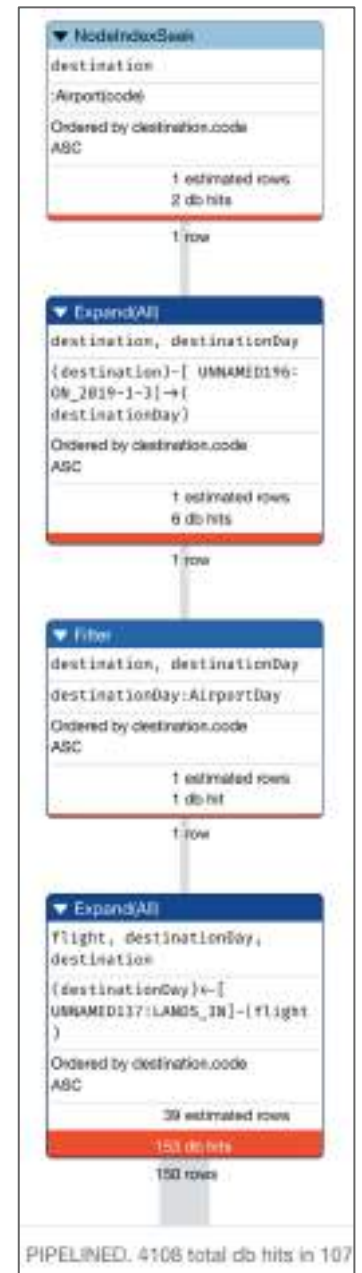
Started streaming 1 records after 1 ms and completed after 75 ms.

# Does the Query Improve?

Rewrite the query, since the model has changed:

## PROFILE

```
MATCH (origin:Airport {code: 'LAS'})-[:`ON_2019-1-3`]->
      (originDay:AirportDay),
      (originDay)-[:ORIGINATES_FROM]-(flight:Flight),
      (flight)-[:LANDS_IN]->(destinationDay),
      (destinationDay:AirportDay)-[:`ON_2019-1-3`]-
      -(destination:Airport {code: 'MDW'})
RETURN flight.date, flight.number, flight.airline,
       flight.departure, flight.arrival
ORDER BY flight.date, flight.departure
```



## Exercise 6: Creating Specific Relationships

In the query edit pane of Neo4j Browser, execute the browser command:

```
:play 4.0-neo4j-modeling-exercises
```

and follow the instructions for Exercise 6.

Note: This exercise has 2 steps. Estimated time to complete: 15 minutes.



# Refactoring Large Graphs

The heap size needs to be adapted to match, or operate in batches.

Increase these values for the server in the [neo4j.conf](#) file to support the need for additional memory:

- `dbms.memory.heap.initial_size=2G` (default [512m](#))
- `dbms.memory.heap.max_size=2G` (default [1G](#))

# Batching the Refactoring Process

1. Tag all the nodes that need to be processed with a temporary label
  - For example *Process* could be used as the temporary label

```
MATCH (f:Flight)
SET f:Process
```

1. Iterate over the subset of nodes flagged with the temporary label
  - This is done using **LIMIT**
    - a. Execute the refactoring code
    - b. Remove the temporary label from the nodes
    - c. Return a count of how many rows were processed
1. The refactoring is done when the count reaches 0



# Example Code for a Batch

```
MATCH (flight:Process)
WITH flight LIMIT 500

MATCH (origin:Airport)<-[:ORIGINATES_FROM]-(flight)-[:LANDS_IN]->(destination:Airport)

MERGE (originAirportDay:AirportDay {airportDayId: origin.code + "_" + flight.date})
ON CREATE SET originAirportDay.date = flight.date

MERGE (destinationAirportDay:AirportDay {airportDayId: destination.code + "_" + flight.date})
ON CREATE SET destinationAirportDay.date = flight.date

MERGE (origin)-[:HAS_DAY]->(originAirportDay)
MERGE (originAirportDay)<-[:ORIGINATES_FROM]-(flight)
MERGE (flight)-[:LANDS_IN]-(destinationAirportDay)
MERGE (destination)-[:HAS_DAY]->(destinationAirportDay)

REMOVE flight:Process
RETURN count(*)
```

# Using apoc.periodic.commit

neo4j> call apoc.help('apoc.periodic')

type	name	text	signature	roles	writes
"procedure"	"apoc.periodic.cancel"	"apoc.periodic.cancel(name) - cancel job with the given name"	"apoc.periodic.cancel(name :: STRING?) :: (name :: STRING?, delay :: INTEGER?, rate :: INTEGER?, done :: BOOLEAN?, cancelled :: BOOLEAN?)"	null	null
"procedure"	"apoc.periodic.commit"	"apoc.periodic.commit(statement,params) - runs the given statement in separate transactions until it returns 0"	"apoc.periodic.commit(statement :: STRING?, params = [] :: MAP?) :: (updates :: INTEGER?, executions :: INTEGER?, runtime :: INTEGER?, batches :: INTEGER?, failedBatches :: INTEGER?, batchErrors :: MAP?, failedCommits :: INTEGER?, commitErrors :: MAP?, wasTerminated :: BOOLEAN?)"	null	null
"procedure"	"apoc.periodic.countdown"	"apoc.periodic.countdown(name',statement,repeat-rate-in-seconds) submit a repeatedly-called background statement until it returns 0"	"apoc.periodic.countdown(name :: STRING?, statement :: STRING?, rate :: INTEGER?) :: (name :: STRING?, delay :: INTEGER?, rate :: INTEGER?, done :: BOOLEAN?, cancelled :: BOOLEAN?)"	null	null
"procedure"	"apoc.periodic.iterate"	"apoc.periodic.iterate(statement returning items', 'statement per item', (batchSize:1000,iterateList:true,parallel:false,params:{},concurrency:50,retries:0)) YIELD batches, total - run the second statement for each item returned by the first statement. Returns number of batches and total processed rows"	"apoc.periodic.iterate(cypherIterate :: STRING?, cypherAction :: STRING?, config :: MAP?) :: (batches :: INTEGER?, total :: INTEGER?, timeTaken :: INTEGER?, committedOperations :: INTEGER?, failedOperations :: INTEGER?, failedBatches :: INTEGER?, retries :: INTEGER?, errorMessages :: MAP?, batch :: MAP?, operations :: MAP?, wasTerminated :: BOOLEAN?, failedParams :: MAP?)"	null	null

Started streaming 6 records after 1 ms and completed after 126 ms.

# Batching with APOC

```
CALL apoc.periodic.commit('
MATCH (flight:Process)
WITH flight LIMIT $limit

MATCH (origin:Airport)<-[:ORIGINATES_FROM]-(flight)-[:LANDS_IN]->(destination:Airport)

MERGE (originAirportDay:AirportDay {airportDayId: origin.code + "_" + flight.date})
ON CREATE SET originAirportDay.date = flight.date

MERGE (destinationAirportDay:AirportDay {airportDayId: destination.code + "_" + flight.date})
ON CREATE SET destinationAirportDay.date = flight.date

MERGE (origin)-[:HAS_DAY]->(originAirportDay)
MERGE (originAirportDay)<-[:ORIGINATES_FROM]-(flight)
MERGE (flight)-[:LANDS_IN]-(destinationAirportDay)
MERGE (destination)-[:HAS_DAY]->(destinationAirportDay)

REMOVE flight:Process
RETURN count(*)

',{limit:500}
)
```

# Result of the Batch Processing

neo4j\$ CALL apoc.periodic.commit(" MATCH (flight:Process) WITH flight LIMIT \$limit MATCH (origin:Airport)-[:ORIGINATES\_FROM]-(flight)-[:LANDS\_IN]->(d...)

updates	executions	runtime	batches	failedBatches	batchErrors	failedCommits	commitErrors	wasTerminated
10000	20	1	21	0	{ }	0	{ }	false

Started streaming 1 records after 1 ms and completed after 1814 ms.

## Exercise 7: Refactoring Large Graphs

In the query edit pane of Neo4j Browser, execute the browser command:

```
:play 4.0-neo4j-modeling-exercises
```

and follow the instructions for Exercise 7.

Note: This exercise has 8 steps. Estimated time to complete: 30 minutes.





# Summary

You should now be able to:

- Create constraints to improve performance of node creation and queries
- Determine if a query can be improved
- Write Cypher code to refactor a data model
- Create indexes that to improve query performance
- Refactor a graph by creating intermediate nodes
- Refactor a graph by specializing relationships
- Perform batch processing when refactoring a large graph



The background of the slide features a pair of hands holding several interlocking puzzle pieces. Overlaid on this is a network graph with various sized nodes and connecting lines. The overall color palette is a mix of soft greens, blues, and yellows.

# Summary

## Implementing Graph Data Models in Neo4j 4.0

# Course Summary

In this course, you have learned how to:

- Implement a simple graph data model
- Import data into an existing graph data model
- Profile query performance against the implemented graph
- Refactor graph data models